

NETWORK PROGRAMMING REPORT

Comparison of Iterative and Concurrent Servers

1. Introduction

In this report, we compare iterative and concurrent servers, analyzing their advantages, disadvantages, real-world applications, and system resource usage.

2. Iterative Server

- Handles one client at a time.
- Processes the client request, then closes the connection before accepting another.
- Simple to implement but slow for handling multiple clients.

Advantages:

- Low resource usage.
- Easy to develop and debug.

Disadvantages:

- Poor scalability (blocks new clients until the current one finishes).
- Not suitable for real-time applications.

Real-World Applications:

- FTP servers (single-user mode).
- Debugging or testing environments.

3. Concurrent Server

- Can handle multiple clients at the same time using threads or processes.
- More efficient for large-scale applications.

Advantages:

- Fast response time.
- Can serve many clients simultaneously.
- Ideal for high-traffic applications.

Disadvantages:

- Higher resource consumption.
- More complex to implement due to thread synchronization issues.

Real-World Applications:

- Web servers (Apache, Nginx).
- Online multiplayer games.
- Real-time stock market systems.

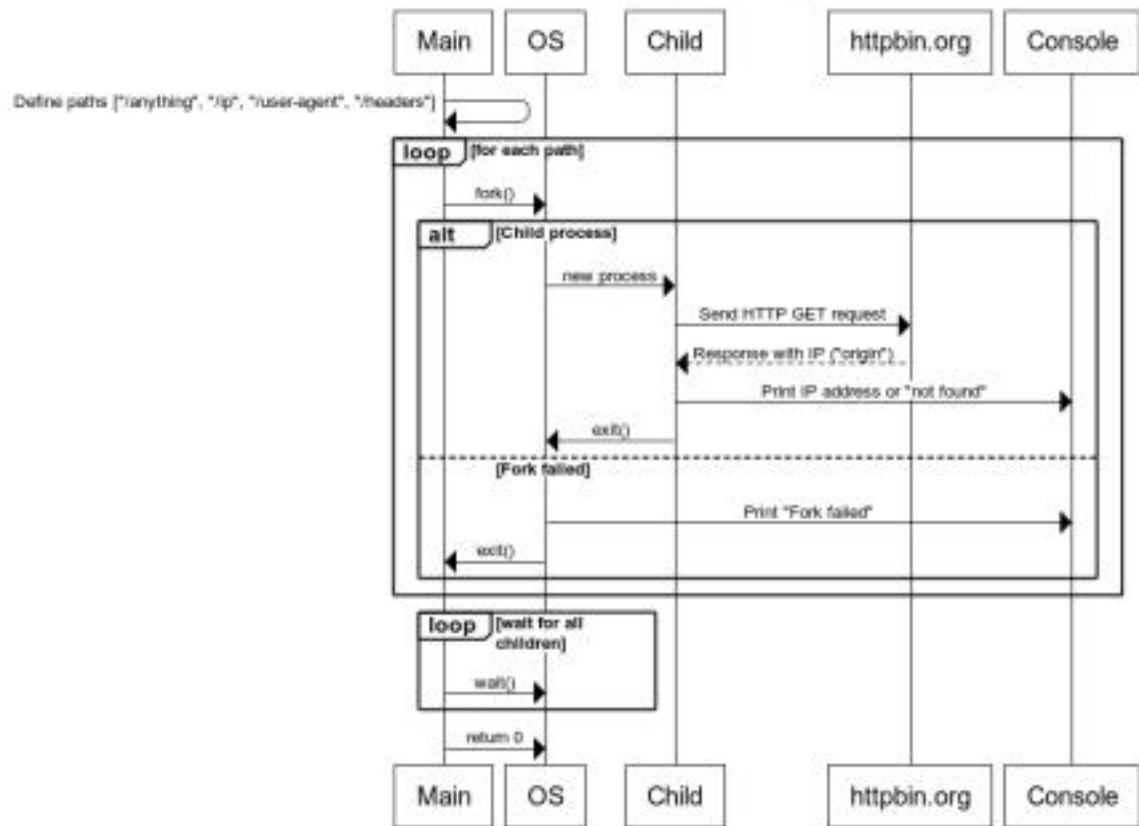
4. System Resource Usage Analysis

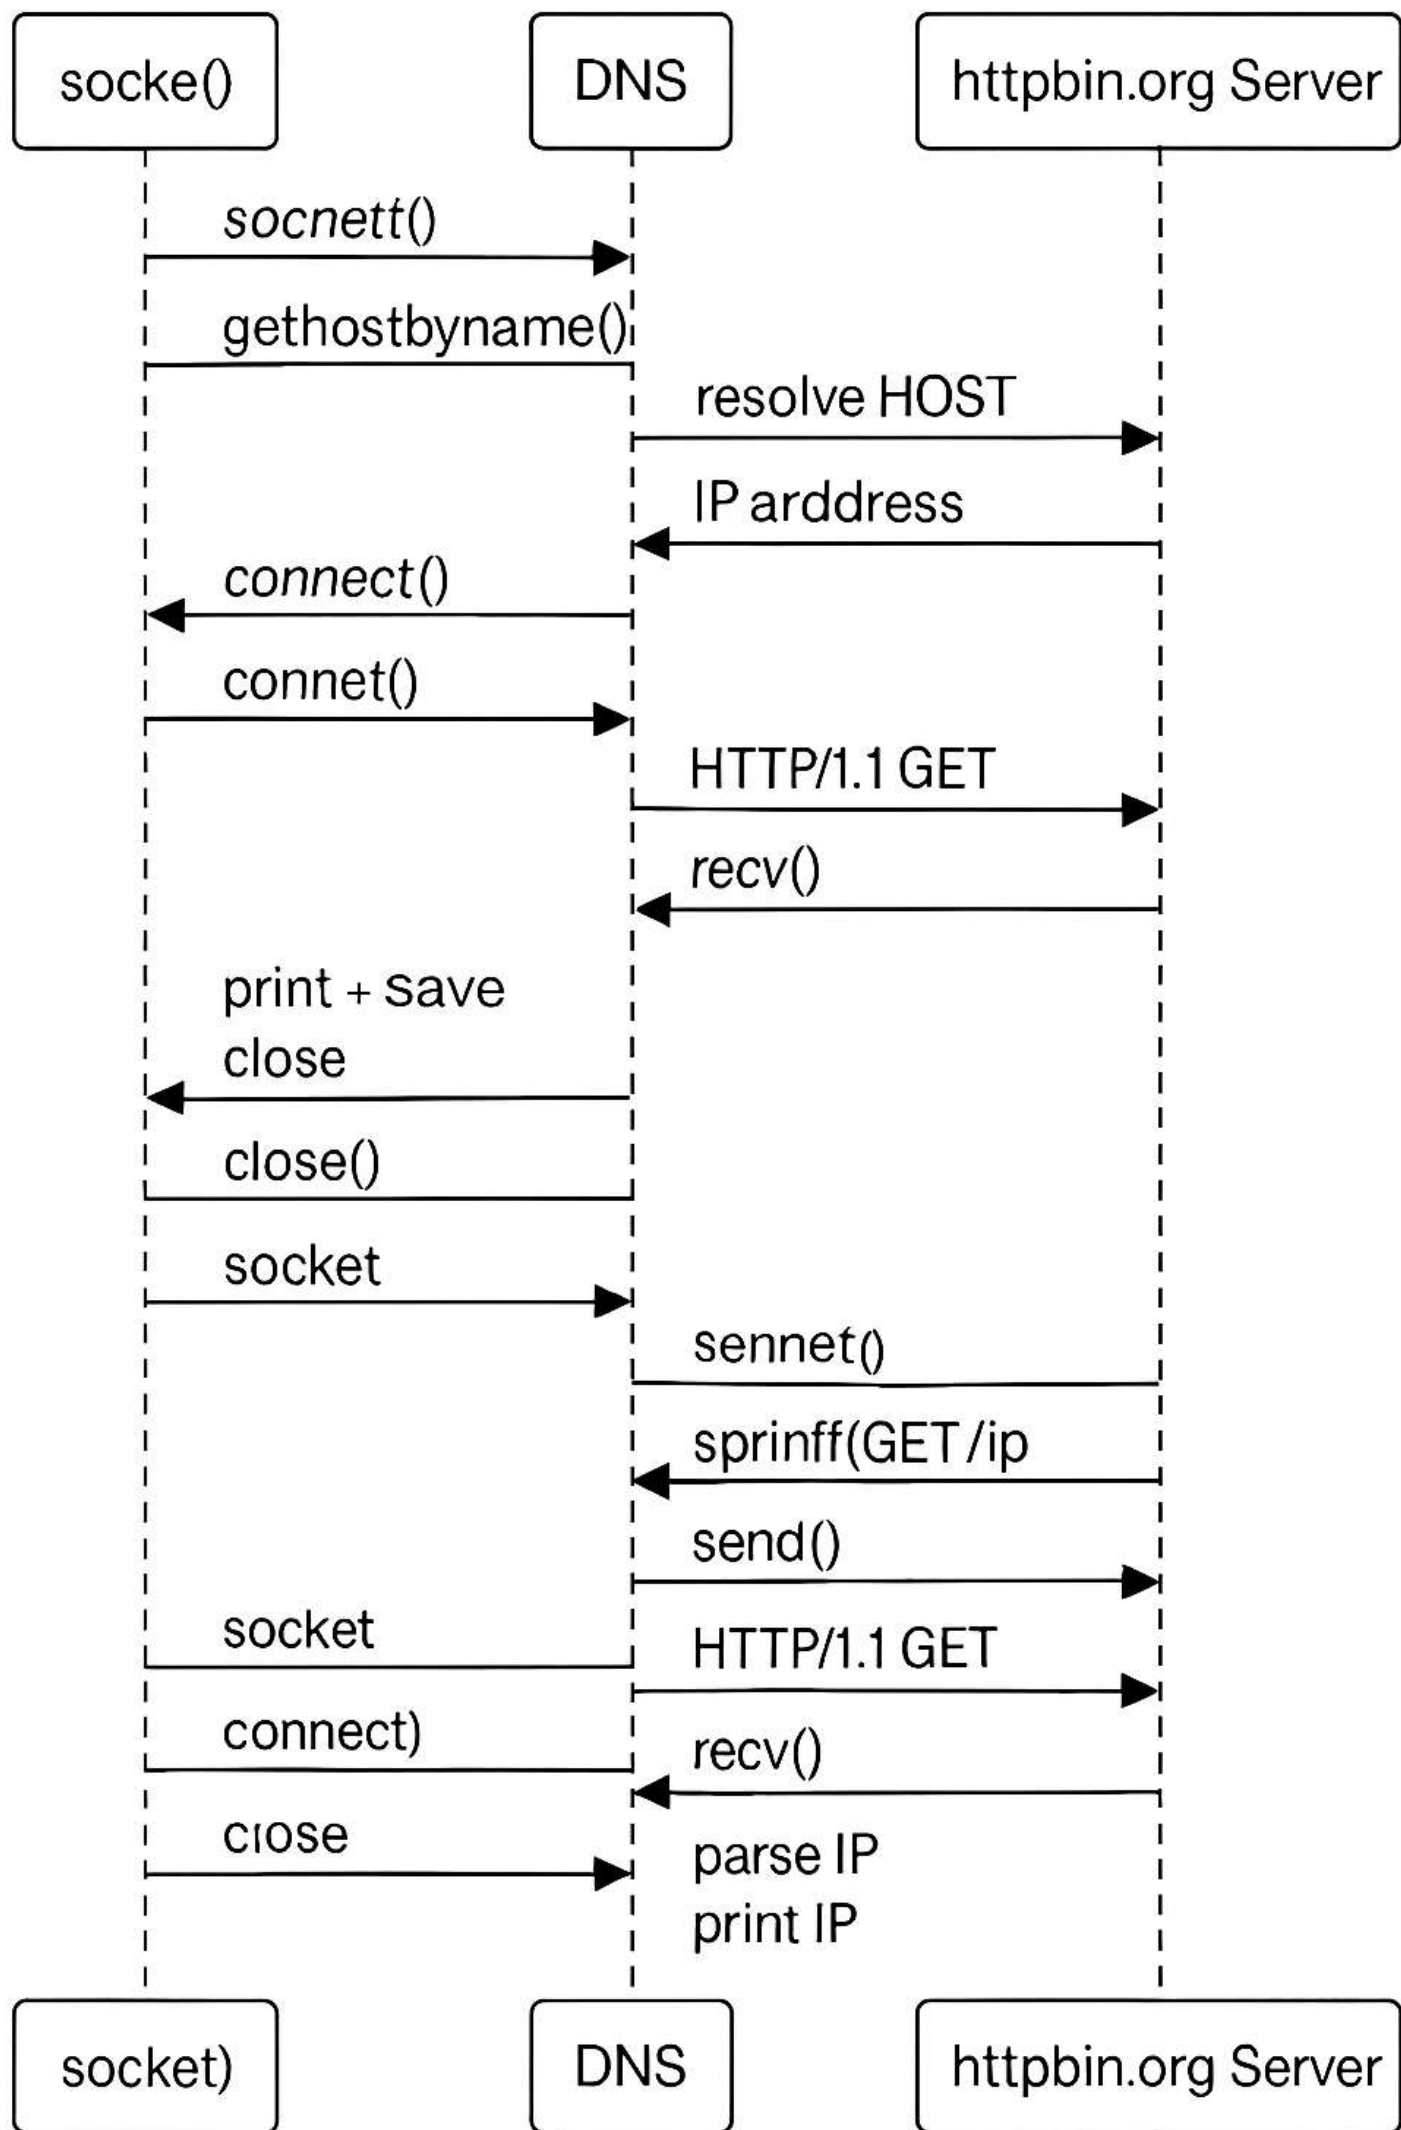
- Iterative servers use minimal resources but are slow under heavy load.
- Concurrent servers handle high loads but consume more CPU and memory.
- System resource usage can be monitored using the command:
ps aux | grep server

5. Conclusion

- Iterative servers are useful for small applications but lack scalability.
- Concurrent servers are preferred for modern applications requiring multiple connections.

Concurrent HTTP Client (Simplified Logic)





```
File Actions Edit View Help

Saving /stream/1 response to stream.html...
HTTP/1.1 200 OK
Date: Tue, 01 Apr 2025 14:36:43 GMT
Content-Type: application/json
Transfer-Encoding: chunked
Connection: close
Server: unicorn/19.9.0
Access-Control-Allow-Origin: *
Access-Control-Allow-Credentials: true

e8
{"url": "http://httpbin.org/stream/1", "args": {}, "headers": {"Host": "httpbin.org", "X-Amzn-Trace-Id": "Root=1-67ebf34b-52df321946/c63ae652c75e0", "User-Agent": "Mozilla/5.0", "Accept": "*/*", "origin": "90.131.46.154", "id": 0}}
0

Fetching client IP address...
Client IP: 90.131.46.154

[mirza@kali] ~/Desktop
$ gcc tcp-client.c -o tcp-client

[mirza@kali] ~/Desktop
$ ./tcp-client

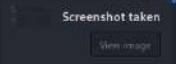
Saving /stream/1 response to stream.html...
HTTP/1.1 200 OK
Date: Tue, 01 Apr 2025 14:24:10 GMT
Content-Type: application/json
Transfer-Encoding: chunked
Connection: close
Server: unicorn/19.9.0
Access-Control-Allow-Origin: *
Access-Control-Allow-Credentials: true

e8
{"url": "http://httpbin.org/stream/1", "args": {}, "headers": {"Host": "httpbin.org", "X-Amzn-Trace-Id": "Root=1-67ebf70a-43b9d0e80f0588d34e35e2d5", "User-Agent": "Mozilla/5.0", "Accept": "*/*", "origin": "90.131.46.154", "id": 0}}
0

Fetching client IP address...
Client IP: 90.131.46.154

[mirza@kali] ~/Desktop
$
```

```
mirza@kali: ~/Desktop
File Actions Edit View Help
Access-Control-Allow-Origin: *
Access-Control-Allow-Credentials: true
{
  "user-agent": "Mozilla/5.0"
}
-- End of /user-agent --
HTTP/1.1 200 OK
Date: Tue, 01 Apr 2025 14:25:23 GMT
Content-Type: application/json
Content-Length: 92
Connection: close
Server: gunicorn/19.9.0
Access-Control-Allow-Origin: *
Access-Control-Allow-Credentials: true
{
  "origin": "90.131.46.154"
}
-- End of /ip --
(mirza@kali)~/Desktop
$
(mirza@kali)~/Desktop
$ gcc concurrentt-client.c -o concurrentt-client
(mirza@kali)~/Desktop
$ ./concurrentt-client
/anything → IP address found: 90.131.46.154
/ip → IP address found: 90.131.46.154
/user-agent → IP address not found
/headers → IP address not found
(mirza@kali)~/Desktop
$ gcc concurrentt-client.c -o concurrentt-client
(mirza@kali)~/Desktop
$ ./concurrentt-client
/ip → IP address found: 90.131.46.154
/anything → IP address found: 90.131.46.154
/user-agent → IP address not found
/headers → IP address not found
(mirza@kali)~/Desktop
$
```



Source Code: Tcp-client.c

```
1: #include <stdio.h>
2: #include <stdlib.h>
3: #include <string.h>
4: #include <unistd.h>
5: #include <netdb.h>
6: #include <arpa/inet.h>
7:
8: #define PORT 80
9: #define HOST "httpbin.org"
10:
11: int main() {
12:     int sock;
13:     struct sockaddr_in server_addr;
14:     struct hostent *server;
15:     char request[1024], response[4096];
16:     int bytes_received;
17:
18:     // -----
19:     // Step 1: GET /stream/1
20:     // -----
21:
22:     // Create socket
23:     sock = socket(AF_INET, SOCK_STREAM, 0);
24:     if (sock < 0) {
25:         perror("Socket error");
26:         return 1;
27:     }
28:
29:     // Get host info
30:     server = gethostbyname(HOST);
31:     if (server == NULL) {
32:         fprintf(stderr, "No such host\n");
33:         return 1;
34:     }
35:
36:     // Set up the server address
37:     memset(&server_addr, 0, sizeof(server_addr));
38:     server_addr.sin_family = AF_INET;
39:     server_addr.sin_port = htons(PORT);
40:     memcpy(&server_addr.sin_addr.s_addr, server->h_addr, server->h_length);
41:
42:     // Connect to the server
43:     if (connect(sock, (struct sockaddr *)&server_addr, sizeof(server_addr)) < 0) {
44:         perror("Connection failed");
45:         return 1;
46:     }
47:
48:     // Send HTTP GET for /stream/1 with headers to avoid 502 error
49:     sprintf(request,
50:         "GET /stream/1 HTTP/1.1\r\n"
51:         "Host: %s\r\n"
```

```

52:         "User-Agent: Mozilla/5.0\r\n"
53:         "Accept: */*\r\n"
54:         "Connection: close\r\n\r\n", HOST);
55: send(sock, request, strlen(request), 0);
56:
57: // Save response to file
58: FILE *file = fopen("stream.html", "w");
59: if (!file) {
60:     perror("File error");
61:     return 1;
62: }
63:
64: printf("Saving /stream/1 response to stream.html...\n");
65: while ((bytes_received = recv(sock, response, sizeof(response) - 1, 0)) > 0) {
66:     response[bytes_received] = '\0';
67:     printf("%s", response); // Print to terminal
68:     fprintf(file, "%s", response); // Save to file
69: }
70:
71: fclose(file);
72: close(sock);
73:
74: // -----
75: // Step 2: GET /ip and extract IP
76: // -----
77:
78: // Create new socket
79: sock = socket(AF_INET, SOCK_STREAM, 0);
80: if (sock < 0) {
81:     perror("Socket error (IP)");
82:     return 1;
83: }
84:
85: // Reconnect to the server
86: if (connect(sock, (struct sockaddr *)&server_addr, sizeof(server_addr)) < 0) {
87:     perror("Connection failed (IP)");
88:     return 1;
89: }
90:
91: // Send HTTP GET for /ip
92: sprintf(request, "GET /ip HTTP/1.1\r\nHost: %s\r\nConnection: close\r\n\r\n", HOST);
93: send(sock, request, strlen(request), 0);
94:
95: printf("\nFetching client IP address...\n");
96:
97: // Receive full response into a buffer
98: char full_response[8192] = "";
99: while ((bytes_received = recv(sock, response, sizeof(response) - 1, 0)) > 0) {
100:     response[bytes_received] = '\0';
101:     strcat(full_response, response);
102: }
103:
104: // Extract the IP address
105: char *origin = strstr(full_response, "\"origin\":");

```

```
106:     if (origin != NULL) {
107:         char ip[64];
108:         if (sscanf(origin, "\"origin\": \"%63[^\"]\"", ip) == 1) {
109:             printf("\nClient IP: %s\n", ip);
110:         } else {
111:             printf("\nCould not parse IP address.\n");
112:         }
113:     } else {
114:         printf("\n'origin' field not found in response.\n");
115:     }
116:
117:     close(sock);
118:     return 0;
119: }
```


Source Code: concurrentt-client.c

```
1: #include <stdio.h>
2: #include <stdlib.h>
3: #include <string.h>
4: #include <unistd.h>
5: #include <netdb.h>
6: #include <arpa/inet.h>
7: #include <sys/wait.h>
8:
9: #define PORT 80
10: #define HOST "httpbin.org"
11:
12: void fetch_ip(const char *path) {
13:     int sock;
14:     struct sockaddr_in server_addr;
15:     struct hostent *server;
16:     char request[1024], response[4096], full_response[8192] = "";
17:     int bytes_received;
18:
19:     sock = socket(AF_INET, SOCK_STREAM, 0);
20:     if (sock < 0) exit(1);
21:
22:     server = gethostbyname(HOST);
23:     if (!server) exit(1);
24:
25:     memset(&server_addr, 0, sizeof(server_addr));
26:     server_addr.sin_family = AF_INET;
27:     server_addr.sin_port = htons(PORT);
28:     memcpy(&server_addr.sin_addr.s_addr, server->h_addr, server->h_length);
29:
30:     if (connect(sock, (struct sockaddr *)&server_addr, sizeof(server_addr)) < 0) exit(1);
31:
32:     sprintf(request,
33:         "GET %s HTTP/1.1\r\n"
34:         "Host: %s\r\n"
35:         "User-Agent: Mozilla/5.0\r\n"
36:         "Connection: close\r\n\r\n", path, HOST);
37:     send(sock, request, strlen(request), 0);
38:
39:     while ((bytes_received = recv(sock, response, sizeof(response) - 1, 0)) > 0) {
40:         response[bytes_received] = '\0';
41:         strcat(full_response, response);
42:     }
43:
44:     // Try to find and extract IP
45:     char *origin = strstr(full_response, "\"origin\":");
46:     if (origin) {
47:         char ip[64];
48:         if (sscanf(origin, "\"origin\": \"%63[^\\\"]\"", ip) == 1) {
49:             printf("%s -> IP address found: %s\n", path, ip);
50:         } else {
51:             printf("%s -> IP address not found\n", path);
```

```
52:     }
53: } else {
54:     printf("%s -> IP address not found\n", path);
55: }
56:
57: close(sock);
58: }
59:
60: int main() {
61:     const char *paths[] = { "/anything", "/ip", "/user-agent", "/headers" };
62:     int num_paths = sizeof(paths) / sizeof(paths[0]);
63:
64:     for (int i = 0; i < num_paths; i++) {
65:         pid_t pid = fork();
66:
67:         if (pid == 0) {
68:             fetch_ip(paths[i]);
69:             exit(0);
70:         } else if (pid < 0) {
71:             perror("Fork failed");
72:             exit(1);
73:         }
74:     }
75:
76:     for (int i = 0; i < num_paths; i++) {
77:         wait(NULL);
78:     }
79:
80:     return 0;
81: }
```